

```
In [1]: #pip install ISLP
```

```
In [2]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import (ModelSpec as MS,
summarize)
```

```
In [3]: from ISLP import confusion_table
from ISLP.models import contrast
from sklearn.discriminant_analysis import \
    (LinearDiscriminantAnalysis as LDA, QuadraticDiscriminantAnalysis as QDA)
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
In [4]: from ISLP import load_data
from ISLP.models import (ModelSpec as MS,
                        summarize,
                        poly)
```

```
In [5]: Auto = load_data("Auto")
```

14a.

```
In [6]: mpg01=(Auto['mpg']>Auto['mpg'].median()).astype(int)
mpg01
```

Out [6]:

	mpg
name	
chevrolet chevelle malibu	0
buick skylark 320	0
plymouth satellite	0
amc rebel sst	0
ford torino	0
...	...
ford mustang gl	1
vw pickup	1
dodge rampage	1
ford ranger	1
chevy s-10	1

392 rows x 1 columns

**dtype:** int64

In [7]: `Auto['mpg01'] = mpg01`

14b.

```
In [8]: #Originally generated by CHATgpt
import matplotlib.pyplot as plt
import seaborn as sns
# Set plot style
sns.set(style="whitegrid")

# Convert 'origin' to categorical with labels
Auto['origin'] = Auto['origin'].replace({1: 'American', 2: 'European', 3: 'Japanese'})

# Plot 1: Cylinders vs mpg01 Boxplot
g1 = sns.boxplot(data=Auto, x='mpg01', y='cylinders', hue='mpg01')
g1.legend_.remove()
plt.title('Cylinders vs mpg01 Boxplot')
plt.show()

# Plot 2: Displacement vs mpg01 Boxplot
g2 = sns.boxplot(data=Auto, x='mpg01', y='displacement', hue='mpg01')
g2.legend_.remove()
plt.title('Displacement vs mpg01 Boxplot')
plt.show()
```

```

# Plot 3: Horsepower vs mpg01 Boxplot
g3 = sns.boxplot(data=Auto, x='mpg01', y='horsepower', hue='mpg01')
g3.legend_.remove()
plt.title('Horsepower vs mpg01 Boxplot')
plt.show()

# Plot 4: Weight vs mpg01 Boxplot
g4 = sns.boxplot(data=Auto, x='mpg01', y='weight', hue='mpg01')
g4.legend_.remove()
plt.title('Weight vs mpg01 Boxplot')
plt.show()


# Plot 5: Acceleration vs mpg01 Boxplot
g5 = sns.boxplot(data=Auto, x='mpg01', y='acceleration', hue='mpg01')
g5.legend_.remove()
plt.title('Acceleration vs mpg01 Boxplot')
plt.show()

# Plot 6: Year vs mpg01 - Boxplot
g6 = sns.boxplot(data=Auto, x='mpg01', y='year', hue='mpg01')
g6.legend_.remove()
plt.title('Year vs mpg01 Boxplot')
plt.show()


# Plot 7: Origin vs mpg01 Bar plot
g7 = sns.histplot(data=Auto, x='origin', hue='mpg01', multiple='fill', shrink=0.8)
plt.gca().yaxis.set_major_formatter(plt.FuncFormatter(lambda y, _: '{:.0%}'))
plt.ylabel('')
plt.title('Origin vs mpg01 Bar plot')
plt.show()

```


Cylinders vs mpg01 Boxplot




Displacement vs mpg01 Boxplot




Horsepower vs mpg01 Boxplot



Weight vs mpg01 Boxplot





Acceleration vs mpg01 Boxplot





Year vs mpg01 Boxplot





Explanation: It looks like all of the other features are useful in predicting mpg01, but I would say cylinders, displacement, weight, and horsepower are more powerful than the others.

14c.

```
In [9]: np.random.seed(123)
train, test = train_test_split(Auto, test_size=0.5, random_state=123)
```

14d.

```
In [10]: # Prepare features and mpg01
X_train = train[['cylinders', 'displacement', 'weight', 'horsepower']]
y_train = train['mpg01']
X_test = test[['cylinders', 'displacement', 'weight', 'horsepower']]
```

```
In [11]: lda = LDA(store_covariance=True)
lda.fit(X_train, y_train)
predicted_lda = lda.predict(X_test)

# Calculate test error
print(np.mean(predicted_lda != test['mpg01']))
```

0.10204081632653061

14e.

```
In [12]: qda = QDA(store_covariance=True)
qda.fit(train[['cylinders', 'displacement', 'weight', 'horsepower']], train[
predicted_qda = qda.predict(X_test)
print(np.mean(predicted_qda != test['mpg01']))
```

0.09693877551020408

14f.

```
In [13]: logit = LogisticRegression(C=1e10, solver='liblinear')
logit.fit(X_train, y_train)
logit_pred = logit.predict(X_test)
print(np.mean(logit_pred != test['mpg01']))
```

0.12755102040816327

14g.

```
In [14]: NB = GaussianNB()
NB.fit(X_train, y_train)
NB_pred=NB.predict(X_test)
print(np.mean(NB_pred != test['mpg01']))
```

0.10204081632653061

14h.

```
In [15]: from sklearn.model_selection import GridSearchCV
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)
print(np.mean(knn_pred != test['mpg01']))
```

0.11734693877551021

The best value of k I could find was 7 since the error starts increasing after k=7.

Test error for k=1: 0.16326530612244897

Test error for k=8: 0.12244897959183673

16.

```
In [16]: Boston = load_data("Boston")
crim01=(Boston['crim']>Boston['crim'].median()).astype(int)
crim01
```

Out[16]:

	crim
0	0
1	0
2	0
3	0
4	0
...	...
501	0
502	0
503	0
504	0
505	0

506 rows x 1 columns

**dtype:** int64


```
In [17]: Boston['crim01'] = crim01
```

```
In [18]: np.random.seed(234)
train, test = train_test_split(Boston, test_size=0.5, random_state=234)
```

```
In [19]: Boston.columns
```

```
Out[19]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
               'ptratio', 'lstat', 'medv', 'crim01'],
              dtype='object')
```

```
In [20]: pd.plotting.scatter_matrix(Boston[['zn', 'indus', 'chas', 'nox', 'rm', 'age',
      'ptratio', 'lstat', 'medv']]);
```



```
In [21]: # Prepare features and crim01
BosAll=['zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax', 'ptratio', 'lstat', 'medv', 'crim01']
X_train = train[BosAll]
y_train = train['crim01']
X_test = test[BosAll]

#logit
logit = LogisticRegression(C=1e10, solver='liblinear')
logit.fit(X_train, y_train)
logit_pred = logit.predict(X_test)
print("Logit:" + str(np.mean(logit_pred != test['crim01'])))

#lda
lda = LDA(store_covariance=True)
lda.fit(X_train, y_train)
predicted_lda = lda.predict(X_test)
print("LDA:" + str(np.mean(predicted_lda != test['crim01'])))

#NB
NB = GaussianNB()
NB.fit(X_train, y_train)
NB_pred=NB.predict(X_test)
print("NB:" + str(np.mean(NB_pred != test['crim01'])))

#KNN
knn = KNeighborsClassifier(n_neighbors=6)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)
print("KNN:" + str(np.mean(knn_pred != test['crim01'])))
```

Logit:0.11462450592885376  
LDA:0.18181818181818182  
NB:0.2015810276679842  
KNN:0.07114624505928854

```
In [22]: # Prepare features and crim01
BosSub=['zn', 'indus', 'rm', 'dis', 'rad', 'tax', 'nox']
X_train = train[BosSub]
y_train = train['crim01']
X_test = test[BosSub]
#logit
logit = LogisticRegression(C=1e10, solver='liblinear')
logit.fit(X_train, y_train)
logit_pred = logit.predict(X_test)
print("Logit:"+str(np.mean(logit_pred != test['crim01'])))
#lda
lda = LDA(store_covariance=True)
lda.fit(X_train, y_train)
predicted_lda = lda.predict(X_test)
print("LDA:"+str(np.mean(predicted_lda != test['crim01'])))
#NB
NB = GaussianNB()
NB.fit(X_train, y_train)
NB_pred=NB.predict(X_test)
print("NB:"+str(np.mean(NB_pred != test['crim01'])))
#KNN
knn = KNeighborsClassifier(n_neighbors=6)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)
print("KNN:"+str(np.mean(knn_pred != test['crim01'])))
```

Logit:0.1857707509881423  
LDA:0.16205533596837945  
NB:0.22134387351778656  
KNN:0.05533596837944664

Findings:

For all predictors, I noticed that the best value of k for KNN is k=6.

KNN also gave the lowest mean error out of the 4 methods for this case.

But since some of the predictors appeared to be correlated with each other, I removed some of them (and created BosSub), which made the logistic and NB test errors to increase, while LDA and KNN test errors decreased (KNN is still the lowest).