

```
In [1]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
from statsmodels.api import OLS
import sklearn.model_selection as skm
import sklearn.linear_model as skl
from sklearn.preprocessing import StandardScaler
from ISLP import load_data
from ISLP.models import ModelSpec as MS
from functools import partial
```

```
In [2]: from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.cross_decomposition import PLSRegression
from ISLP.models import \
    (Stepwise,
     sklearn_selected,
     sklearn_selection_path)
from l0bnb import fit_path
```

```
In [3]: from sklearn.exceptions import ConvergenceWarning
import warnings
warnings.filterwarnings("ignore", category=ConvergenceWarning)
warnings.filterwarnings("ignore", category=UserWarning)
```

```
In [4]: Hitters = load_data('Hitters')
np.isnan(Hitters['Salary']).sum()
```

Out[4]: 59

```
In [5]: Hitters = Hitters.dropna()
Hitters.shape
```

Out[5]: (263, 20)

```
In [6]: def nCp(sigma2, estimator, X, Y):
    "Negative Cp statistic"
    n, p = X.shape
    Yhat = estimator.predict(X)
    RSS = np.sum((Y - Yhat)**2)
    return -(RSS + 2 * p * sigma2) / n
```

```
In [7]: design = MS(Hitters.columns.drop('Salary')).fit(Hitters)
Y = np.array(Hitters['Salary'])
X = design.transform(Hitters)
sigma2 = OLS(Y, X).fit().scale
```

```
In [8]: neg_Cp = partial(nCp, sigma2)
```

```
In [9]: strategy = Stepwise.first_peak(design,  
                                     direction='forward',  
                                     max_terms=len(design.terms))
```

```
In [10]: hitters_MSE = sklearn_selected(OLS,  
                                     strategy)  
hitters_MSE.fit(Hitters, Y)  
hitters_MSE.selected_state_
```

```
Out[10]: ('Assists',  
          'AtBat',  
          'CAatBat',  
          'CHits',  
          'CHmRun',  
          'CRBI',  
          'CRuns',  
          'CWalks',  
          'Division',  
          'Errors',  
          'Hits',  
          'HmRun',  
          'League',  
          'NewLeague',  
          'PutOuts',  
          'RBI',  
          'Runs',  
          'Walks',  
          'Years')
```

```
In [11]: hitters_Cp = sklearn_selected(OLS,  
                                     strategy,  
                                     scoring=neg_Cp)  
hitters_Cp.fit(Hitters, Y)  
hitters_Cp.selected_state_
```

```
Out[11]: ('Assists',  
          'AtBat',  
          'CAatBat',  
          'CRBI',  
          'CRuns',  
          'CWalks',  
          'Division',  
          'Hits',  
          'PutOuts',  
          'Walks')
```

```
In [12]: strategy = Stepwise.fixed_steps(design,  
                                     len(design.terms),  
                                     direction='forward')  
full_path = sklearn_selection_path(OLS, strategy)
```


```
In [13]: full_path.fit(Hitters, Y)  
Yhat_in = full_path.predict(Hitters)
```

```
Yhat_in.shape
```

```
Out[13]: (263, 20)
```

```
In [14]: mse_fig, ax = subplots(figsize=(8, 8))
insample_mse = ((Yhat_in - Y[:, None]) ** 2).mean(0)
n_steps = insample_mse.shape[0]
ax.plot(np.arange(n_steps),
        insample_mse,
        'k', # color black
        label='In-sample')
ax.set_ylabel('MSE',
              fontsize=20)
ax.set_xlabel('# steps of forward stepwise',
              fontsize=20)
ax.set_xticks(np.arange(n_steps)[:,2])
ax.legend()
ax.set_ylim([50000, 250000])
```

```
Out[14]: (50000.0, 250000.0)
```



```
In [15]: K = 5
kfold = skm.KFold(K,
                  random_state=0,
                  shuffle=True)
Yhat_cv = skm.cross_val_predict(full_path,
                                Hitters,
                                Y,
                                cv=kfold)


Yhat_cv.shape
```

Out[15]: (263, 20)

```
In [16]: cv_mse = []
for train_idx, test_idx in kfold.split(Y):
    errors = (Yhat_cv[test_idx] - Y[test_idx, None]) ** 2
    cv_mse.append(errors.mean(0)) # column means
cv_mse = np.array(cv_mse).T
cv_mse.shape
```

Out[16]: (20, 5)

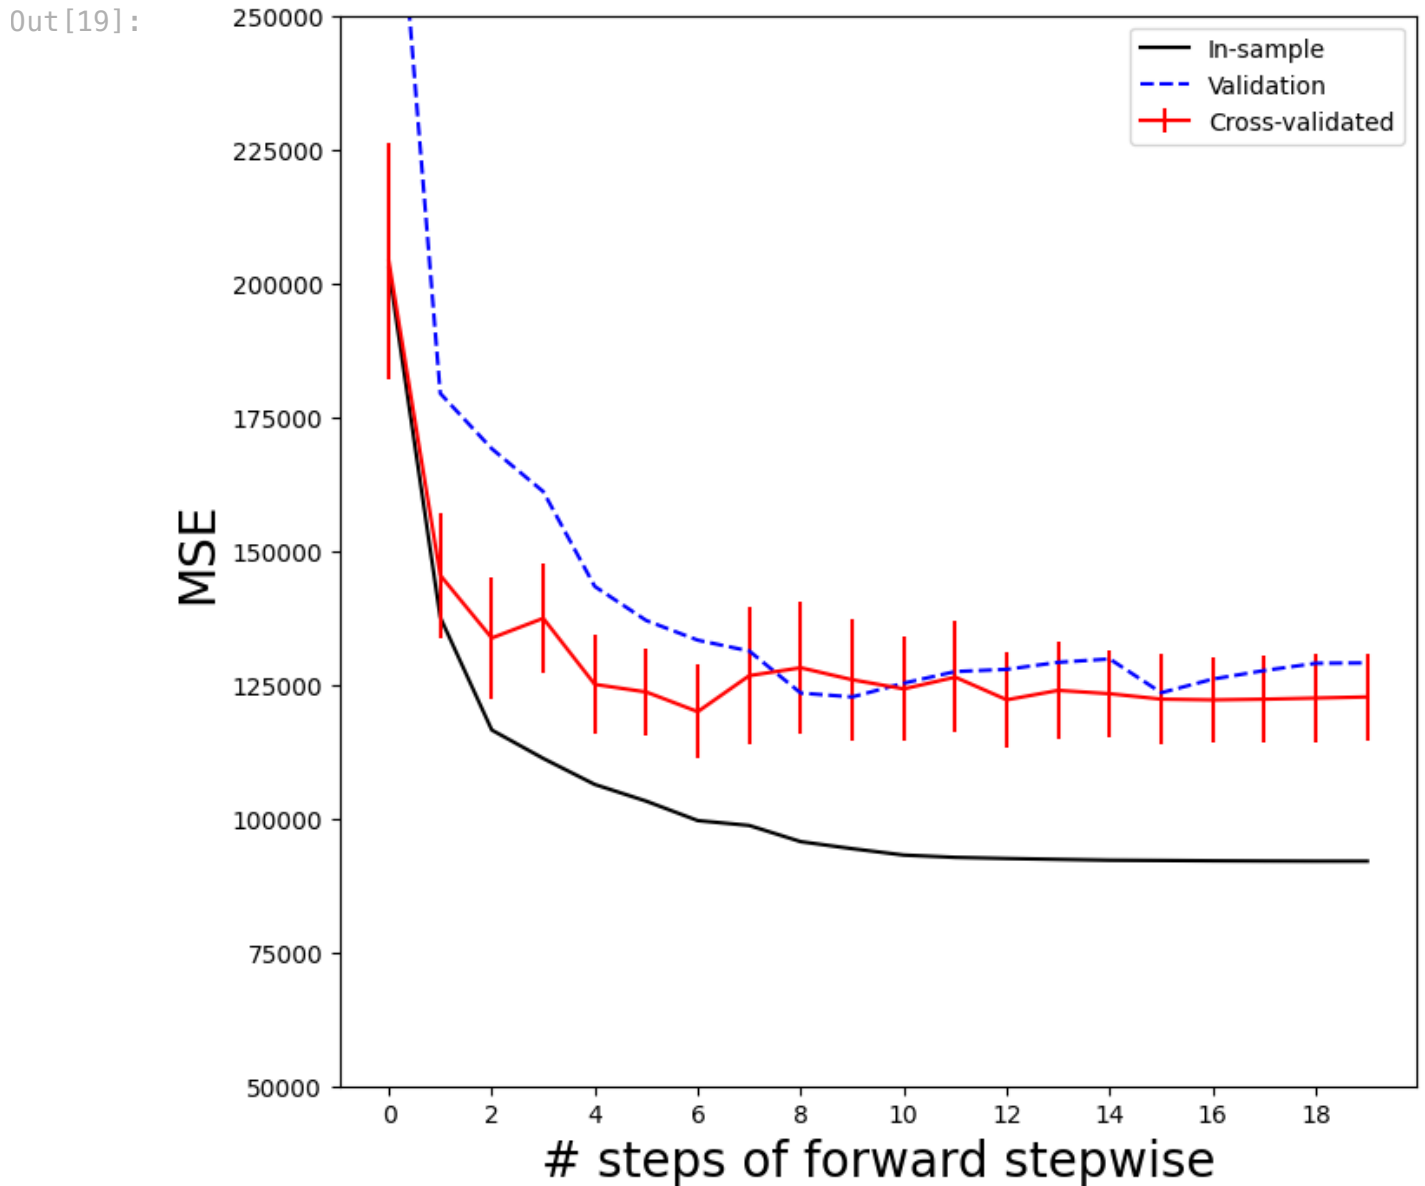
```
In [17]: ax.errorbar(np.arange(n_steps),
                    cv_mse.mean(1),
                    cv_mse.std(1) / np.sqrt(K),
                    label='Cross-validated',
                    c='r') # color red
ax.set_ylim([50000, 250000])
ax.legend()
mse_fig
```



```
In [18]: validation = skm.ShuffleSplit(n_splits=1,
                                       test_size=0.2,
                                       random_state=0)
for train_idx, test_idx in validation.split(Y):
    full_path.fit(Hitters.iloc[train_idx],
                  Y[train_idx])
    Yhat_val = full_path.predict(Hitters.iloc[test_idx])
    errors = (Yhat_val - Y[test_idx, None]) ** 2
```

```
validation_mse = errors.mean(0)
```

```
In [19]: ax.plot(np.arange(n_steps),  
                validation_mse,  
                'b--', # color blue , broken line  
                label='Validation')  
ax.set_xticks(np.arange(n_steps)[:2])  
ax.set_ylim([50000, 250000])  
ax.legend()  
mse_fig
```



```
In [20]: D = design.fit_transform(Hitters)  
D = D.drop('intercept', axis=1)  
X = np.asarray(D)
```

```
In [21]: path = fit_path(X,  
                        Y,  
                        max_nonzeros=X.shape[1])
```

Preprocessing Data.

BnB Started.

```
Iteration: 1. Number of non-zeros: 1
Iteration: 2. Number of non-zeros: 2
Iteration: 3. Number of non-zeros: 2
Iteration: 4. Number of non-zeros: 2
Iteration: 5. Number of non-zeros: 3
Iteration: 6. Number of non-zeros: 3
Iteration: 7. Number of non-zeros: 4
Iteration: 8. Number of non-zeros: 9
Iteration: 9. Number of non-zeros: 9
Iteration: 10. Number of non-zeros: 9
Iteration: 11. Number of non-zeros: 9
Iteration: 12. Number of non-zeros: 9
Iteration: 13. Number of non-zeros: 9
Iteration: 14. Number of non-zeros: 9
Iteration: 15. Number of non-zeros: 9
Iteration: 16. Number of non-zeros: 9
Iteration: 17. Number of non-zeros: 9
Iteration: 18. Number of non-zeros: 17
Iteration: 19. Number of non-zeros: 19
```

In [22]: `path[3]`

```
Out[22]: {'B': array([0.          , 3.25484367, 0.          , 0.          , 0.          ,
                    0.          , 0.          , 0.          , 0.          ,
                    0.          , 0.67775265, 0.          , 0.          , 0.          ,
                    0.          , 0.          , 0.          , 0.          ]),
          'B0': -38.98216739555551,
          'lambda_0': 0.011416248027450178,
          'M': 0.5829861733382012,
          'Time_exceeded': False}
```

```
In [23]: Xs = X - X.mean(0)[None, :]
X_scale = X.std(0)
Xs = Xs / X_scale[None, :]
lambdas = 10 ** np.linspace(8, -2, 100) / Y.std()
soln_array = skl.ElasticNet.path(Xs,
                                Y,
                                l1_ratio=0.,
                                alphas=lambdas)[1]

soln_array.shape
```

Out[23]: (19, 100)

```
In [24]: soln_path = pd.DataFrame(soln_array.T,
                                columns=D.columns,
                                index=-np.log(lambdas))
soln_path.index.name = 'negative log(lambda)'
soln_path
```

Out[24]:

	AtBat	Hits	HmRun	Runs	RBI	Walk
negative log(lambda)						
-12.310855	0.000800	0.000889	0.000695	0.000851	0.000911	0.00090
-12.078271	0.001010	0.001122	0.000878	0.001074	0.001150	0.00113
-11.845686	0.001274	0.001416	0.001107	0.001355	0.001451	0.00143
-11.613102	0.001608	0.001787	0.001397	0.001710	0.001831	0.00180
-11.380518	0.002029	0.002255	0.001763	0.002158	0.002310	0.00228
...	...	...	...	...	...	...
9.784658	-290.823989	336.929968	37.322686	-59.748520	-26.507086	134.85591
10.017243	-290.879272	337.113713	37.431373	-59.916820	-26.606957	134.90054
10.249827	-290.923382	337.260446	37.518064	-60.051166	-26.686604	134.93613
10.482412	-290.958537	337.377455	37.587122	-60.158256	-26.750044	134.96447
10.714996	-290.986528	337.470648	37.642077	-60.243522	-26.800522	134.98702

100 rows x 19 columns


In [25]:

```
path_fig, ax = subplots(figsize=(8, 8))
soln_path.plot(ax=ax, legend=False)
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Standardized coefficients', fontsize=20)
ax.legend(loc='upper left')
```

```
<>:3: SyntaxWarning: invalid escape sequence '\l'
<>:3: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/2938439178.
py:3: SyntaxWarning: invalid escape sequence '\l'
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
```

Out[25]: <matplotlib.legend.Legend at 0x304386690>





```
In [26]: beta_hat = soln_path.loc[soln_path.index[39]]
         lambdas[39], beta_hat
```

```
Out[26]: (25.53538897200662,
          AtBat          5.433750
          Hits           6.223582
          HmRun          4.585498
          Runs           5.880855
          RBI            6.195921
          Walks          6.277975
          Years          5.299767
          CAtBat         7.147501
          CHits          7.539495
          CHmRun         7.182344
          CRuns          7.728649
          CRBI           7.790702
          CWalks         6.592901
          League[N]      0.042445
          Division[W]    -3.107159
          PutOuts        4.605263
          Assists        0.378371
          Errors         -0.135196
          NewLeague[N]   0.150323
          Name: -3.240065292879872, dtype: float64)
```


```
In [27]: np.linalg.norm(beta_hat)
```

```
Out[27]: 24.17061720144378
```

```
In [28]: beta_hat = soln_path.loc[soln_path.index[59]]
         lambdas[59], np.linalg.norm(beta_hat)
```

```
Out[28]: (0.24374766133488554, 160.42371017725839)
```

```
In [29]: ridge = skl.ElasticNet(alpha=lambdas[59], l1_ratio=0)
         scaler = StandardScaler(with_mean=True, with_std=True)
         pipe = Pipeline(steps=[('scaler', scaler), ('ridge', ridge)])
         pipe.fit(X, Y)
```



```
In [30]: np.linalg.norm(ridge.coef_)
```

```
Out[30]: 160.4237101772591
```

```
In [31]: validation = skm.ShuffleSplit(n_splits=1,
                                       test_size=0.5,
                                       random_state=0)
```

```

ridge.alpha = 0.01
results = skm.cross_validate(ridge,
                              X,
                              Y,
                              scoring='neg_mean_squared_error',
                              cv=validation)

results['test_score']

```

Out[31]: array([134214.00419204])

```

In [32]: ridge.alpha = 1e10
results = skm.cross_validate(ridge,
                              X,
                              Y,
                              scoring='neg_mean_squared_error',
                              cv=validation)

results['test_score']

```


Out[32]: array([231788.32155285])

```

In [33]: param_grid = {'ridge__alpha': lambdas}
grid = skm.GridSearchCV(pipe,
                        param_grid,
                        cv=validation,
                        scoring='neg_mean_squared_error')

grid.fit(X, Y)
grid.best_params_['ridge__alpha']
grid.best_estimator_

```



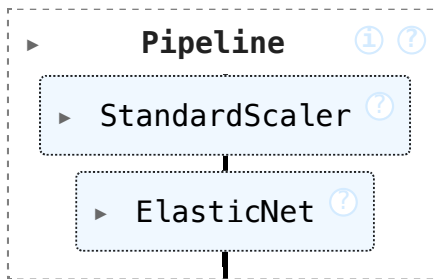
```

In [34]: grid = skm.GridSearchCV(pipe,
                                param_grid,
                                cv=kfold,
                                scoring='neg_mean_squared_error')

grid.fit(X, Y)
grid.best_params_['ridge__alpha']
grid.best_estimator_

```

Out[34]:




In [35]:

```
ridge_fig, ax = subplots(figsize=(8, 8))
ax.errorbar(-np.log(lambdas),
            -grid.cv_results_['mean_test_score'],
            yerr=grid.cv_results_['std_test_score'] / np.sqrt(K))
ax.set_ylim([50000, 250000])
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Cross-validated MSE', fontsize=20)
```

```
<>:6: SyntaxWarning: invalid escape sequence '\l'
<>:6: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/252384649.p
y:6: SyntaxWarning: invalid escape sequence '\l'
    ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
```

Out[35]: Text(0, 0.5, 'Cross-validated MSE')



```
In [36]: grid_r2 = skm.GridSearchCV(pipe,
                                     param_grid,
                                     cv=kfold)
grid_r2.fit(X, Y)
```

```
Out[36]:
```

- **GridSearchCV** ⓘ ?
  - **best\_estimator\_: Pipeline**
    - StandardScaler ⓘ
    - ElasticNet ⓘ

```
In [37]: r2_fig, ax = subplots(figsize=(8, 8))
ax.errorbar(-np.log(lambdas),
```

```

        grid_r2.cv_results_['mean_test_score'],
        yerr=grid_r2.cv_results_['std_test_score'] / np.sqrt(K)
    )
    ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
    ax.set_ylabel('Cross-validated  $R^2$ ', fontsize=20)


```

```

<>:6: SyntaxWarning: invalid escape sequence '\l'
<>:6: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/4088780906.
py:6: SyntaxWarning: invalid escape sequence '\l'
    ax.set_xlabel('$-\log(\lambda)$', fontsize=20)

```

Out[37]: Text(0, 0.5, 'Cross-validated  $R^2$ ')




```

In [38]: ridgeCV = skl.ElasticNetCV(alphas=lambdas,
                                     l1_ratio=0,
                                     cv=kfold)
pipeCV = Pipeline(steps=[('scaler', scaler),

```

```
pipeCV.fit(X, Y)
('ridge', ridgeCV)]
```

Out[38]:




In [39]:

```
tuned_ridge = pipeCV.named_steps['ridge']
ridgeCV_fig, ax = subplots(figsize=(8, 8))
ax.errorbar(-np.log(lambdas),
            tuned_ridge.mse_path_.mean(1),
            yerr=tuned_ridge.mse_path_.std(1) / np.sqrt(K))
ax.axvline(-np.log(tuned_ridge.alpha_), c='k', ls='--')
ax.set_ylim([50000, 250000])
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Cross-validated MSE', fontsize=20)
```

```
<>:8: SyntaxWarning: invalid escape sequence '\l'
<>:8: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/35476348.p
y:8: SyntaxWarning: invalid escape sequence '\l'
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
```

Out[39]: Text(0, 0.5, 'Cross-validated MSE')



```
In [40]: np.min(tuned_ridge.mse_path_.mean(1))
```

```
Out[40]: 115526.70630987917
```

```
In [41]: tuned_ridge.coef_
```

```
Out[41]: array([-222.80877051,  238.77246614,    3.21103754,  -2.93050845,
                3.64888723,  108.90953869,  -50.81896152, -105.15731984,
                122.00714801,   57.1859509 ,  210.35170348,  118.05683748,
               -150.21959435,   30.36634231,  -61.62459095,   77.73832472,
                40.07350744,  -25.02151514,  -13.68429544])
```

```
In [42]: outer_valid = skm.ShuffleSplit(n_splits=1,
                                         test_size=0.25,
                                         random_state=1)

inner_cv = skm.KFold(n_splits=5,
                    shuffle=True,
                    random_state=2)
```



```
ridgeCV = skl.ElasticNetCV(alphas=lambdas,
                           l1_ratio=0,
                           cv=inner_cv)
pipeCV = Pipeline(steps=[('scaler', scaler),
                          ('ridge', ridgeCV)])
```

```
In [43]: results = skm.cross_validate(pipeCV,
                                     X,
                                     Y,
                                     cv=outer_valid,
                                     scoring='neg_mean_squared_error')

- results['test_score']
```

Out[43]: array([132393.84003227])

```
In [44]: lassoCV = skl.ElasticNetCV(n_alphas=100,
                                   l1_ratio=1,
                                   cv=kfold)

pipeCV = Pipeline(steps=[('scaler', scaler),
                          ('lasso', lassoCV)])

pipeCV.fit(X, Y)
tuned_lasso = pipeCV.named_steps['lasso']
tuned_lasso.alpha_
```

Out[44]: 3.1472370031649866


```
In [45]: lambdas, soln_array = skl.Lasso.path(Xs,
                                              Y,
                                              l1_ratio=1,
                                              n_alphas=100)[:2]

soln_path = pd.DataFrame(soln_array.T,
                        columns=D.columns,
                        index=-np.log(lambdas))
```

```
In [46]: path_fig, ax = subplots(figsize=(8, 8))
soln_path.plot(ax=ax, legend=False)
ax.legend(loc='upper left')
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Standardized coefficients', fontsize=20)
```

```
<>:4: SyntaxWarning: invalid escape sequence '\l'
<>:4: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/1325931816.
py:4: SyntaxWarning: invalid escape sequence '\l'
    ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
```

Out[46]: Text(0, 0.5, 'Standardized coefficients')



```
In [47]: np.min(tuned_lasso.mse_path_.mean(1))
```

```
Out[47]: 114690.73118253727
```


```
In [48]: lassoCV_fig, ax = subplots(figsize=(8, 8))
ax.errorbar(-np.log(tuned_lasso.alphas_),
            tuned_lasso.mse_path_.mean(1),
            yerr=tuned_lasso.mse_path_.std(1) / np.sqrt(K))
ax.axvline(-np.log(tuned_lasso.alpha_), c='k', ls='--')
ax.set_ylim([50000, 250000])
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Cross-validated MSE', fontsize=20)
```

```

<>:7: SyntaxWarning: invalid escape sequence '\l'
<>:7: SyntaxWarning: invalid escape sequence '\l'
/var/folders/97/23ltc4v96g31pp78_gyv6dvm0000gn/T/ipykernel_10519/3583132959.
py:7: SyntaxWarning: invalid escape sequence '\l'
    ax.set_xlabel('$-\log(\lambda)$', fontsize=20)

```

Out[48]: Text(0, 0.5, 'Cross-validated MSE')



In [49]: `tuned_lasso.coef_`

```

Out[49]: array([-210.01008773, 243.4550306, 0., 0.,
                0., 97.69397357, -41.52283116, -0.,
                0., 39.62298193, 205.75273856, 124.55456561,
                -126.29986768, 15.70262427, -59.50157967, 75.24590036,
                21.62698014, -12.04423675, -0.])

```

In [50]: `pca = PCA(n_components=2)`  
`linreg = skl.LinearRegression()`  
`pipe = Pipeline([('pca', pca),`

```

        ('linreg', linreg)])
pipe.fit(X, Y)
pipe.named_steps['linreg'].coef_

```

Out[50]: array([0.09846131, 0.4758765 ])

```

In [51]: pipe = Pipeline([('scaler', scaler),
                           ('pca', pca),
                           ('linreg', linreg)])
pipe.fit(X, Y)
pipe.named_steps['linreg'].coef_

```

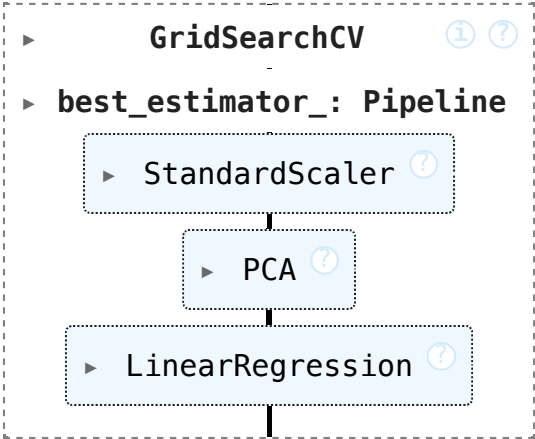
Out[51]: array([106.36859204, 21.60350456])

```

In [52]: param_grid = {'pca__n_components': range(1, 20)}
grid = skm.GridSearchCV(pipe,
                        param_grid,
                        cv=kfold,
                        scoring='neg_mean_squared_error')
grid.fit(X, Y)

```

Out[52]:



```

GridSearchCV
├── best_estimator_: Pipeline
│   ├── StandardScaler
│   │   └── PCA
│   │       └── LinearRegression

```

```

In [53]: pcr_fig, ax = subplots(figsize=(8, 8))
n_comp = param_grid['pca__n_components']
ax.errorbar(n_comp,
            -grid.cv_results_['mean_test_score'],
            grid.cv_results_['std_test_score'] / np.sqrt(K))
ax.set_ylabel('Cross-validated MSE', fontsize=20)
ax.set_xlabel('# principal components', fontsize=20)
ax.set_xticks(n_comp[::2])
ax.set_ylim([50000, 250000])

```

Out[53]: (50000.0, 250000.0)



Out[56]:

▼ PLSRegression ⓘ ?  
PLSRegression()

In [57]:

```
param_grid = {'n_components': range(1, 20)}  
grid = skm.GridSearchCV(pls,  
                        param_grid,  
                        cv=kfold,  
                        scoring='neg_mean_squared_error')  
grid.fit(X, Y)
```

Out[57]:

► GridSearchCV ⓘ ?  
 ► best\_estimator\_: PLSRegression  
 ► PLSRegression ⓘ

In [58]:

```
pls_fig, ax = subplots(figsize=(8, 8))  
n_comp = param_grid['n_components']  
ax.errorbar(n_comp,  
            -grid.cv_results_['mean_test_score'],  
            grid.cv_results_['std_test_score'] / np.sqrt(K))  
ax.set_ylabel('Cross-validated MSE', fontsize=20)  
ax.set_xlabel('# principal components', fontsize=20)  
ax.set_xticks(n_comp[::2])  
ax.set_ylim([50000, 250000])
```

Out[58]:

(50000.0, 250000.0)

