

STAT 702 - Homework 3

Noah Javadi

2025-03-09

```
##Setup
```

```
#install.packages('ISLR2','boot','glmnet','pls')  
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.3.3
```

```
library(boot)  
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.3
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.3.3
```

```
##
```

```
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
## loadings
```

```
##Problem 5-9
```

```
#Problem 9a  
u <- mean(Boston$medv)  
u
```

```
## [1] 22.53281
```

```
#Problem 9b
```

```
se1 <- sd(Boston$medv)/sqrt(nrow(Boston))  
se1
```

```
## [1] 0.4088611
```

```
#Problem 9c - The bootstrap estimation is very close to the manual calculation from b
```

```
se.fn <- function(Boston, index) {  
  x <- Boston$medv[index]  
  y <- length(index)  
  
  se <- sd(x)/sqrt(y)  
  se  
}  
se.fn(Boston,1:500)
```

```
## [1] 0.4130334
```

```
boot(Boston,se.fn,R=1000)
```

```
##  
## ORDINARY NONPARAMETRIC BOOTSTRAP  
##  
##  
## Call:  
## boot(data = Boston, statistic = se.fn, R = 1000)  
##  
##  
## Bootstrap Statistics :  
##      original      bias      std. error  
## t1* 0.4088611 -0.001018028  0.01657769
```

```
#Problem 9d - The results from the manual calculation and t.test are very close as well
```

```
CI.u <- c(u - 1.96*boot(Boston,se.fn,R=1000)[[1]],u + 1.96*boot(Boston,se.fn,R=1000)[[1]])  
CI.u
```

```
## [1] 21.73144 23.33417
```

```
t.test(Boston$medv)
```

```
##  
## One Sample t-test  
##  
## data: Boston$medv  
## t = 55.111, df = 505, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 0  
## 95 percent confidence interval:  
## 21.72953 23.33608  
## sample estimates:  
## mean of x  
## 22.53281
```

```
#Problem 9e  
median(Boston$medv)
```

```
## [1] 21.2
```

```
#Problem 9f -  
se.fn(Boston,1:quantile(Boston$medv,probs = c(0.5)))
```

```
## [1] 1.375661
```

```
#Problem 9g  
u0.1 <- mean(quantile(Boston$medv),probs = c(0.1))  
u0.1
```

```
## [1] 23.645
```

```
#Problem 9h -  
se.fn(Boston,1:quantile(Boston$medv,probs = c(0.1)))
```

```
## [1] 2.080688
```

```
##Problem 6-9
```

```
#Problem 9a  
College.Train <- College[1:388,]  
College.Test <- College[389:777,]  
  
College.Trainy <- sample(1:nrow(College),nrow(College)/2)  
College.Testy <- (-College.Trainy)  
Apps.test <- College$Apps[College.Testy]
```

```
#Problem 9b  
model.train <- lm(Apps ~ .,data = College.Train)  
summary(model.train)
```

```
##  
## Call:  
## lm(formula = Apps ~ ., data = College.Train)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -2721.8  -337.1   -30.6   253.5  6341.8   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -7.681e+02  4.895e+02  -1.569  0.11748      
## PrivateYes  -4.834e+02  1.748e+02  -2.766  0.00596 **   
## Accept       1.202e+00  7.875e-02  15.267 < 2e-16 ***  
## Enroll       8.369e-02  2.780e-01   0.301  0.76358      
## Top10perc    4.062e+01  6.950e+00   5.845 1.11e-08 ***
```

```
## Top25perc    -1.353e+01  5.692e+00  -2.377  0.01795 *
## F.Undergrad   3.144e-02  4.338e-02   0.725  0.46902
## P.Undergrad   6.651e-03  5.587e-02   0.119  0.90530
## Outstate     -2.514e-02  2.306e-02  -1.090  0.27638
## Room.Board    1.872e-01  5.813e-02   3.220  0.00139 **
## Books         -1.958e-01  2.599e-01  -0.753  0.45166
## Personal      1.062e-01  8.226e-02   1.291  0.19756
## PhD           7.152e-01  6.115e+00   0.117  0.90695
## Terminal     -1.088e+01  6.780e+00  -1.604  0.10956
## S.F.Ratio     1.168e+01  1.508e+01   0.774  0.43915
## perc.alumni  -6.553e+00  5.098e+00  -1.285  0.19944
## Expend        8.440e-02  1.415e-02   5.963  5.78e-09 ***
## Grad.Rate     7.445e+00  3.515e+00   2.118  0.03483 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 884.6 on 370 degrees of freedom
## Multiple R-squared:  0.9072, Adjusted R-squared:  0.9029
## F-statistic: 212.7 on 17 and 370 DF,  p-value: < 2.2e-16
```

```
#Problem 9c
grid <- 10^seq(10, -2, length = 388)
ridge.train <- glmnet(College.Train[,c(3:18)],College.Train$Apps,alpha=0,thresh = 1e-12)

ridge.mod <- glmnet(model.matrix(Apps ~ .,College.Train),College.Trainy,alpha = 0,lambda = grid)

cv.out <- cv.glmnet(model.matrix(Apps ~ .,College.Train),College.Trainy,alpha=0)

bestlam <- cv.out$lambda.min

ridge.pred <- predict(ridge.mod, s=bestlam, newx = model.matrix(Apps ~ .,College.Train))
#ridge.pred

mean((ridge.pred - College.Trainy)^2)
```

```
## [1] 50486.42
```

```
#Problem 9d
lasso.train <- glmnet(College.Train[,c(1,3:18)],College.Train$Apps,alpha=1)

lasso.mod <- glmnet(model.matrix(Apps ~ .,College.Train),College.Trainy,alpha = 1,lambda = grid)

cv.out <- cv.glmnet(model.matrix(Apps ~ .,College.Train),College.Trainy,alpha=1)

bestlam <- cv.out$lambda.min

lasso.pred <- predict(lasso.mod, s=bestlam, newx = model.matrix(Apps ~ .,College.Train))
#lasso.pred

lasso.pred <- predict(lasso.mod, s=bestlam, type = 'coefficients')[1:19,]
#lasso.pred

mean((lasso.pred - College.Trainy)^2)
```

```
## Warning in lasso.pred - College.Trainy: longer object length is not a multiple
## of shorter object length
```

```
## [1] 200053.7
```

```
#Problem 9e - M = 10
```

```
pcr.fit <- pcr(Apps ~ ., data = College, scale = TRUE, validation = "CV")
summary(pcr.fit)
```

```
## Data:      X dimension: 777 17
## Y dimension: 777 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3873    3836    2024    2030    1748    1587    1579
## adjCV           3873    3837    2022    2030    1651    1579    1577
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1570    1540    1498    1491    1495    1497    1502
## adjCV        1571    1534    1495    1489    1492    1494    1499
##      14 comps 15 comps 16 comps 17 comps
## CV          1503    1443    1172    1131
## adjCV        1500    1424    1165    1125
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          31.670   57.30   64.30   69.90   75.39   80.38   83.99   87.40
## Apps       2.316   73.06   73.07   82.08   84.08   84.11   84.32   85.18
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          90.50   92.91   95.01   96.81   97.9    98.75   99.36
## Apps       85.88   86.06   86.06   86.10   86.1    86.13   90.32
##      16 comps 17 comps
## X          99.84   100.00
## Apps       92.52   92.92
```

```
pcr.pred <- predict(pcr.fit, College.Test, ncomp = 10)
#pcr.pred
mean((pcr.pred - Apps.test)^2)
```

```
## [1] 31777007
```

```
#Problem 9f - M = 10
```

```
pls.fit <- plsrf(Apps ~ ., data = College, subset = College.Trainy, scale = TRUE, validation = "CV")
summary(pls.fit)
```

```
## Data:      X dimension: 388 17
## Y dimension: 388 1
## Fit method: kernelpls
```

```
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3631    1607    1473    1279    1234    1154    1103
## adjCV           3631    1604    1470    1275    1222    1133    1094
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1082    1080    1080    1076    1075    1075    1074
## adjCV        1077    1076    1075    1071    1070    1070    1069
##      14 comps 15 comps 16 comps 17 comps
## CV          1073    1073    1073    1073
## adjCV        1068    1068    1068    1068
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          26.75   53.38   62.95   65.66   67.76   72.06   76.33   80.62
## Apps       81.18   84.72   88.85   90.58   92.08   92.36   92.46   92.50
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          82.68   84.69   87.67   90.74   92.67   95.26   96.89
## Apps       92.58   92.63   92.64   92.64   92.65   92.65   92.65
##      16 comps 17 comps
## X          99.13  100.00
## Apps       92.65   92.65
```

```
pls.pred <- predict(pls.fit,College.Test,ncomp = 10)
#pcr.pred
mean((pls.pred - Apps.test)^2)
```

```
## [1] 33543289
```

#Problem 9g - It does not seem like we are able to accurately predict the number of college applications received based on the data provided. The MSE squared is quite large for the prediction vs test data. The test errors do seem similar between the 5 approaches.

##Problem 6-11

```
#Problem 11a
#Linear
Boston.lm <- lm(crim ~ .,data = Boston)
summary(Boston.lm)
```


```
##
## Call:
## lm(formula = crim ~ ., data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.534 -2.248 -0.348  1.087  73.923
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```


```
## (Intercept) 13.7783938 7.0818258 1.946 0.052271 .
## zn          0.0457100 0.0187903 2.433 0.015344 *
## indus      -0.0583501 0.0836351 -0.698 0.485709
## chas       -0.8253776 1.1833963 -0.697 0.485841
## nox        -9.9575865 5.2898242 -1.882 0.060370 .
## rm         0.6289107 0.6070924 1.036 0.300738
## age        -0.0008483 0.0179482 -0.047 0.962323
## dis        -1.0122467 0.2824676 -3.584 0.000373 ***
## rad         0.6124653 0.0875358 6.997 8.59e-12 ***
## tax        -0.0037756 0.0051723 -0.730 0.465757
## ptratio    -0.3040728 0.1863598 -1.632 0.103393
## lstat       0.1388006 0.0757213 1.833 0.067398 .
## medv       -0.2200564 0.0598240 -3.678 0.000261 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.46 on 493 degrees of freedom
## Multiple R-squared:  0.4493, Adjusted R-squared:  0.4359
## F-statistic: 33.52 on 12 and 493 DF,  p-value: < 2.2e-16
```


```
Boston.red.lm <- lm(crim ~ dis + rad + medv + zn, data = Boston)
summary(Boston.red.lm)
```


```
##
## Call:
## lm(formula = crim ~ dis + rad + medv + zn, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.459 -1.960 -0.331  0.857 74.718
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.26548    1.34674   3.910 0.000105 ***
## dis         -0.72291    0.20254  -3.569 0.000393 ***
## rad          0.50021    0.04044  12.370 < 2e-16 ***
## medv        -0.19122    0.03566  -5.362 1.26e-07 ***
## zn           0.05487    0.01735   3.163 0.001658 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.5 on 501 degrees of freedom
## Multiple R-squared:  0.4335, Adjusted R-squared:  0.429
## F-statistic: 95.84 on 4 and 501 DF,  p-value: < 2.2e-16
```

```
plot(Boston.red.lm)
```









```
#Ridge
Boston.x <- model.matrix(crim ~ ., Boston)[, -1]
Boston.y <- Boston$crim
Boston.lambda <- 10^seq(10, -2, length = 100)

Boston.train <- sample(1:nrow(Boston), nrow(Boston)/2)
Boston.test <- (-Boston.train)
Boston.ytest <- Boston.y[Boston.test]

Boston.ridge.mod <- glmnet(Boston.x, Boston.y, alpha = 0, lambda = Boston.lambda)
predict(Boston.ridge.mod, s=0, type = 'coefficients')
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 13.4473681349
## zn          0.0452630596
## indus       -0.0605142483
## chas        -0.8223648009
## nox         -9.7788562633
## rm          0.6281672462
## age        -0.0008916473
## dis        -1.0038844850
## rad         0.6051810433
## tax        -0.0034139067
## ptratio    -0.2990248708
## lstat      0.1401134551
```

```
## medv          -0.2178870802
```

```
Boston.ridge.mod <- glmnet(Boston.x[Boston.train,], Boston.y[Boston.train], alpha = 0, lambda = Boston.
Boston1.cv.out <- cv.glmnet(Boston.x[Boston.train,],Boston.y[Boston.train],alpha=0)
summary(Boston.ridge.mod)
```

```
##           Length Class      Mode
## a0          100  -none-   numeric
## beta        1200 dgCMatrix S4
## df           100  -none-   numeric
## dim           2  -none-   numeric
## lambda       100  -none-   numeric
## dev.ratio    100  -none-   numeric
## nulldev       1  -none-   numeric
## npasses       1  -none-   numeric
## jerr          1  -none-   numeric
## offset        1  -none-   logical
## call          5  -none-    call
## nobs          1  -none-   numeric
```

```
Boston1.bestlam <- Boston1.cv.out$lambda.min
```

```
Boston.ridge.pred <- predict(Boston.ridge.mod,s=Boston1.bestlam,newx = Boston.x[Boston.test,])
Boston.s.pred <- predict(Boston.lm,newdata = Boston[Boston.test,])
summary(Boston.ridge.pred)
```

```
##           s1
## Min.      :-3.14605
## 1st Qu.: 0.08169
## Median : 1.21295
## Mean     : 3.36367
## 3rd Qu.: 3.82836
## Max.     :14.84202
```

```
summary(Boston$crim)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00632 0.08204 0.25651 3.61352 3.67708 88.97620
```

```
mean((Boston.s.pred - Boston.ytest)^2)
```


```
## [1] 47.13302
```

```
mean((Boston.ridge.pred - Boston.ytest)^2)
```


```
## [1] 50.61527
```

```
#Lasso
Boston.lasso.mod <- glmnet(Boston.x[Boston.train,],Boston.y[Boston.train],alpha = 1,lambda = grid)
plot(Boston.lasso.mod)
```

```
## Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):  
## collapsing to unique 'x' values
```



```
Boston2.cv.out <- cv.glmnet(Boston.x[Boston.train,],Boston.y[Boston.train],alpha=1)  
plot(cv.out)
```



```
Boston2.bestlam <- Boston2.cv.out$lambda.min
bestlam
```

```
## [1] 25.07386
```

```
Boston.lasso.pred <- predict(Boston.lasso.mod, s=Boston2.bestlam, newx = Boston.x[Boston.test,])
#Boston.lasso.pred

mean((Boston.lasso.pred - Boston.ytest)^2)
```


```
## [1] 50.00845
```

```
#PCR
Boston.pcr.fit <- pcr(crim ~ ., data = Boston, scale = TRUE, validation = "CV")
summary(Boston.pcr.fit)
```

```
## Data:      X dimension: 506 12
## Y dimension: 506 1
## Fit method: svdpc
## Number of components considered: 12
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
```

```
## CV          8.61    7.249    7.243    6.836    6.81    6.766    6.755
## adjCV       8.61    7.247    7.241    6.833    6.81    6.764    6.752
##           7 comps  8 comps  9 comps 10 comps 11 comps 12 comps
## CV         6.631    6.649    6.644    6.636    6.578    6.493
## adjCV      6.626    6.645    6.639    6.631    6.571    6.487
##
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          49.93    63.64    72.94    80.21    86.83    90.26    92.79    94.99
## crim       29.39    29.55    37.39    37.85    38.85    39.23    41.73    41.82
##           9 comps 10 comps 11 comps 12 comps
## X          96.78    98.33    99.48    100.00
## crim       42.12    42.43    43.58    44.93
```

```
validationplot(Boston.pcr.fit, val.type = "MSEP")
```



```
Boston.pcr.pred <- predict(Boston.pcr.fit, Boston.x[Boston.test,], ncomp = 12)
#Boston.pcr.pred

Boston.pcr.fit1 <- pcr(Boston.y ~ Boston.x, scale = TRUE, ncomp = 12)
summary(Boston.pcr.fit1$fitted.values)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -8.0696 -0.2981  1.4915  3.6135  8.4130 17.7592
```

```
mean((Boston.pcr.pred - Boston.ytest)^2)
```

```
## [1] 47.13302
```

#Problem 11b - The PCR model seems to more accurately fit the data and is slightly improved from the linear model. The PCR model also better describes the data based on plots of the Boston crime rate per capita.

#Problem 11c - The PCR model does contain all components of the data set because that is where we achieve the most variability accounted for in Boston crime rate per capita.